

Lab 16. Playbooks and Roles

In this lab, we will cover the following topics:

- Understanding the playbook framework
- Understanding roles---the playbook organizer
- Using conditions in your code
- Repeating tasks with loops
- Grouping tasks using blocks
- Configuring play execution via strategies
- Using [ansible-pull]

Lab Environment

All lab file are present at below path. Run following command in the terminal first before running commands in the lab:

```
cd ~/Desktop/gitlab-ci-ansible-course/Lab_16
```

Understanding the playbook framework

Although YAML format is easy to read and write, it is very pedantic when it comes to spacing. For example, you cannot use tabs to set indentation even though on the screen a tab and four spaces might look identical---in YAML, they are not. We recommend that you adopt an editor with YAML support to aid you in writing your playbooks if you are doing this for the first time, perhaps Vim, Visual Studio Code, or Eclipse, as these will help you to ensure that your indentation is correct. To test the playbooks we develop in this lab, we will reuse a variant of an inventory created in Lab 3 (unless stated otherwise):

```
[frontends]
frt01.example.com https_port=8443
frt02.example.com http_proxy=proxy.example.com

[frontends:vars]
ntp_server=ntp.frt.example.com
proxy=proxy.frt.example.com

[apps]
app01.example.com
app02.example.com

[webapp:children]
frontends
apps

[webapp:vars]
proxy_server=proxy.webapp.example.com
health_check_retry=3
health_check_interal=60
```

Let's dive right in and get started writing a playbook.

1. Create a simple playbook to run on the hosts in the [frontends] host group defined in our inventory file. We can set the user that will access the hosts using the [remote_user] directive in the playbook as demonstrated

in the following (you can also use the `--user` switch on the command line, but as this lab is about playbook development, we'll ignore that for now):

```
---
- hosts: frontends
  remote_user: root

  tasks:
  - name: simple connection test
    ping:
      remote_user: root
```

2. Add another task below the first to run the `[shell]` module (that will, in turn, run the `[ls]` command on the remote hosts). We'll also add the `[ignore_errors]` directive to this task to ensure that our playbook doesn't fail if the `[ls]` command fails (for example, if the directory we're trying to list doesn't exist). Be careful with the indentation and ensure it matches that of the first part of the file:

```
- name: run a simple command
  shell: /bin/ls -al /nonexistent
  ignore_errors: True
```

Let's see how our newly created playbook behaves when we run it:

```
$ ansible-playbook -i hosts myplaybook.yaml

PLAY [frontends] *****

TASK [Gathering Facts] *****
ok: [frr02.example.com]
ok: [frr01.example.com]

TASK [simple connection test] *****
ok: [frr01.example.com]
ok: [frr02.example.com]

TASK [run a simple command] *****
fatal: [frr02.example.com]: FAILED! => {"changed": true, "cmd": "/bin/ls -al
/nonexistent", "delta": "0:00:00.015687", "end": "2020-04-10 16:37:56.895520", "msg":
"non-zero return code", "rc": 2, "start": "2020-04-10 16:37:56.879833", "stderr":
"/bin/ls: cannot access /nonexistent: No such file or directory", "stderr_lines":
["/bin/ls: cannot access /nonexistent: No such file or directory"], "stdout": "",
"stdout_lines": []}
...ignoring
fatal: [frr01.example.com]: FAILED! => {"changed": true, "cmd": "/bin/ls -al
/nonexistent", "delta": "0:00:00.012160", "end": "2020-04-10 16:37:56.930058", "msg":
"non-zero return code", "rc": 2, "start": "2020-04-10 16:37:56.917898", "stderr":
"/bin/ls: cannot access /nonexistent: No such file or directory", "stderr_lines":
["/bin/ls: cannot access /nonexistent: No such file or directory"], "stdout": "",
"stdout_lines": []}
...ignoring

PLAY RECAP *****
```

```
frr01.example.com : ok=3 changed=1 unreachable=0 failed=0 skipped=0 rescued=0
ignored=1
frr02.example.com : ok=3 changed=1 unreachable=0 failed=0 skipped=0 rescued=0
ignored=1
```

From the output of the playbook run, you can see that our two tasks were executed in the order in which they were specified. We can see that the `[ls]` command failed because we tried to list a directory that did not exist, but the playbook did not register any `[failed]` tasks because we set `[ignore_errors]` to `[true]` for this task (and only this task).

A handler is a special type of task that is run as a result of a `[notify]`. However, unlike Ansible playbook tasks, which are performed in sequence, handlers are all grouped together and run at the very end of the play. Also, they can be notified more than once but will only be run once regardless, again preventing needless service restarts. Consider the following playbook:

Change Apache Port

Let's change the default Apache port using template as port 80 is already in use by lab environment.

```
---
- name: Handler demo 1
  hosts: frr01.example.com
  gather_facts: no
  become: yes

  tasks:
    - name: Update Apache configuration
      template:
        src: template.j2
        dest: /etc/apache2/ports.conf
      notify: Restart Apache

  handlers:
    - name: Restart Apache
      service:
        name: apache2
        state: restarted
```

Note: Make sure that `apache2` is installed before running the playbook: `service apache2 status`

You can install `apache2` by running `apt-get install -y apache2`. It can also be installed by running following playbook:

```
$ cd ~/Desktop/gitlab-ci-ansible-course/Lab_14
$ ansible-playbook update-apache-version.yml
```

If we run this task a first time, we will see the following results:

```
$ ansible-playbook -i hosts handlers1.yml

PLAY [Handler demo 1] *****

TASK [Update Apache configuration] *****
changed: [frr01.example.com]
```

```
RUNNING HANDLER [Restart Apache] *****
changed: [frt01.example.com]

PLAY RECAP *****
frt01.example.com : ok=2 changed=2 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

ProTip

Open <http://localhost:81> in Mirdori browser to verify that apache is running now.

Notice how the handler was run at the end, as the configuration file was updated. However, if we run this playbook a second time without making any changes to the template or configuration file, we will see something like this:

```
$ ansible-playbook -i hosts handlers1.yml

PLAY [Handler demo 1] *****

TASK [Update Apache configuration] *****
ok: [frt01.example.com]

PLAY RECAP *****
frt01.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

This time, the handler was not called as the result from the configuration task as OK. All handlers should have a globally unique name so that the notify action can call the correct handler. You could also call multiple handlers by setting a common name for using the [listen] directive---this way, you can call either the handler [name] or the [listen] string---as demonstrated in the following example:

```
---
- name: Handler demo 1
  hosts: frt01.example.com
  gather_facts: no
  become: yes

  handlers:
    - name: restart ntp
      service:
        name: ntp
        state: restarted
      listen: "restart all services"
    - name: restart apache
      service:
        name: apache2
        state: restarted
      listen: "restart all services"

  tasks:
    - name: restart all services
      command: echo "this task will restart all services"
      notify: "restart all services"
```

We only have one task in the playbook, but when we run it, both handlers are called. Also, remember that we said earlier that [command] was among a set of modules that were a special case because they can't detect whether a change has occurred---as a result, they always return the [changed] value, and so, in this demo playbook, the handlers will always be notified:

```
$ ansible-playbook -i hosts handlers2.yml

PLAY [Handler demo 1] *****

TASK [restart all services] *****
changed: [frt01.example.com]

RUNNING HANDLER [restart ntp] *****
changed: [frt01.example.com]

RUNNING HANDLER [restart apache] *****
changed: [frt01.example.com]

PLAY RECAP *****
frt01.example.com : ok=3 changed=3 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

These are some of the fundamentals that you need to know to start writing your own playbooks. With these under your belt, let's run through a comparison of ad hoc commands and playbooks in the next section.

Comparing playbooks and ad hoc tasks

Let's develop a practical example---suppose you want to install Apache. There are a number of steps involved even if the default configuration is sufficient (which is unlikely, but for now, we'll keep the example simple). If you were to perform the basic installation by hand, you would need to install the package, open up the firewall, and ensure the service is running (and runs at boot time).

To perform these commands in the shell, you might do the following:

```
$ sudo apt install apache2
$ sudo service apache2 start
$ sudo service apache2 status
```

Note: Start will fail because port 80 is already in use by lab environment. make sure tha
/etc/apache2/ports.conf is updated; port 80 with port 81 .

Now, for each of these commands, there is an equivalent ad hoc Ansible command that you could run. We won't go through all of them here in the interests of space; however, let's say you want to restart the Apache service---in this case, you could run an ad hoc command similar to the following (again, we will perform it only on one host for conciseness):

```
$ ansible -i hosts frt01* -m service -a "name=apache2 state=restarted"
```

When run successfully, you will see pages of shell output containing all of the variable data returned from running the service module in this way. A snippet of this is shown in the following for you to check yours against---the key thing being that the command resulted in the [changed] status, meaning that it ran successfully and that the service was indeed restarted:

```
frt01.example.com | CHANGED => {
  "ansible_facts": {
    "discovered_interpreter_python": "/usr/bin/python"
  },
  "changed": true,
  "name": "apache2",
  "state": "started",
```

A playbook is a far more valuable way to approach this---not only will it perform all of the steps in one go, but it will also give you a record of how it was done for you to refer to later on. There are multiple ways to do this, but consider the following as an example:

```
---
- name: Install Apache
  hosts: frt01.example.com
  gather_facts: no
  become: yes

  tasks:
    - name: Install Apache package
      apt:
        name: apache2
        state: latest
    - name: Restart and enable the service
      service:
        name: apache2
        state: restarted
        enabled: yes
```

Now, when you run this, you should see that all of our installation requirements have been completed by one fairly simple and easy to read playbook. There is a new concept here, loops, which we haven't covered yet, but don't worry, we will cover this later in this lab:

```
$ ansible-playbook -i hosts installapache.yml

PLAY [Install Apache] *****

TASK [Install Apache package] *****
changed: [frt01.example.com]

TASK [Restart and enable the service] *****
changed: [frt01.example.com]

PLAY RECAP *****
frt01.example.com : ok=2 changed=3 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

As you can see, this is far better for capturing what was actually done and documenting it in a format that someone else could easily pick up. Even though we will cover loops later on in the course, it's fairly easy to see from the preceding how they might be working. With this set out, let's proceed in the next section to look in more detail at a couple of terms we have used several times to ensure you are clear on their meanings: **plays** and **tasks**.

Defining plays and tasks

Suppose we want to write a single playbook to configure both the frontend servers and the application servers. We could use two separate playbooks to configure the front end and application servers, but this risks fragmenting your code and making it difficult to organize. However, front end servers and application servers are going to be (by their very nature) fundamentally different and so are unlikely to be configured with the same set of tasks.

The solution to this problem is to create a single playbook with two plays in it. The start of each play can be identified by the line at the lowest indentation (that is, zero spaces in front of it). Let's get started with building up our playbook:

1. Add the first play to the playbook and define some simple tasks to set up the Apache server on the front end, as shown here:

```
---
- name: Play 1 - configure the frontend servers
  hosts: frontends
  become: yes

  tasks:
    - name: Install the Apache package
      apt:
        name: apache2
        state: latest
    - name: Start the Apache server
      service:
        name: apache2
        state: started
```

2. Immediately below this, in the same file, add the second play to configure the application tier servers:

```
- name: Play 2 - configure the application servers
  hosts: apps
  become: true

  tasks:
    - name: Install NTP
      apt:
        name: ntp
        state: latest
    - name: Start the NTP server
      service:
        name: ntp
        state: started
```

Now, you have two plays: one to install web servers in the [frontends] group and one to install application servers in the [apps] group, all combined into one simple playbook.

When we run this playbook, we'll see the two plays performed sequentially, in the order they appear in the playbook. Note the presence of the [PLAY] keyword, which denotes the start of each play:

```
$ ansible-playbook -i hosts playandtask.yml
```

```

PLAY [Play 1 - configure the frontend servers] *****

TASK [Gathering Facts] *****
changed: [frr02.example.com]
changed: [frr01.example.com]

TASK [Install the Apache package] *****
changed: [frr01.example.com]
changed: [frr02.example.com]

TASK [Start the Apache server] *****
changed: [frr01.example.com]
changed: [frr02.example.com]

PLAY [Play 2 - configure the application servers] *****

TASK [Gathering Facts] *****
changed: [app01.example.com]
changed: [app02.example.com]

TASK [Install NTP] *****
changed: [app02.example.com]
changed: [app01.example.com]

TASK [Start the NTP server] *****
changed: [app02.example.com]
changed: [app01.example.com]

PLAY RECAP *****
app01.example.com : ok=3 changed=2 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
app02.example.com : ok=3 changed=2 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frr01.example.com : ok=3 changed=2 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frr02.example.com : ok=3 changed=2 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0

```

There we have it---one playbook, yet two distinct plays operating on different sets of hosts from the provided inventories. This is very powerful, especially when combined with roles (which will be covered later in this course). Of course, you can have just one play in your playbook---you don't have to have multiple ones, but it is important to be able to develop multi-play playbooks as you will almost certainly find them useful as your environment gets more complex.

Understanding roles -- the playbook organizer

The process of creating roles is in fact very simple---Ansible will (by default) look within the same directory as you are running your playbook from for a [roles/] directory, and in here, you will create one subdirectory for each role. The role name is derived from the subdirectory name---there is no need to create complex metadata or anything else---it really is that simple. Within each subdirectory goes a fixed directory structure that tells Ansible what the tasks, default variables, handlers, and so on are for each role.

The [roles/] directory is not the only place Ansible will look for roles---this is the first directory it will look in, but it will then look in [/etc/ansible/roles] for any additional roles. This can be further customized through the Ansible configuration file.

Let's explore this in a little more detail. Consider the following directory structure:

```
site.yml
frontends.yml
dbservers.yml
roles/
  installapache/
    tasks/
    handlers/
    templates/
    vars/
    defaults/
  installntp/
    tasks/
    meta/
```

For the examples we will develop in this part of this lab, we will need an inventory, so let's reuse the inventory we used in the previous section (included in the following for convenience):

```
[frontends]
frt01.example.com https_port=8443
frt02.example.com http_proxy=proxy.example.com

[frontends:vars]
ntp_server=ntp.frt.example.com
proxy=proxy.frt.example.com

[apps]
app01.example.com
app02.example.com

[webapp:children]
frontends
apps

[webapp:vars]
proxy_server=proxy.webapp.example.com
health_check_retry=3
health_check_interval=60
```

Let's get started with some practical exercises to help you to learn how to create and work with roles. We'll start by creating a role called [installapache], which will handle the Apache installation process we looked at in the previous section. However, here, we will expand it to cover the installation of Apache on both CentOS and Ubuntu. This is good practice, especially if you are looking to submit your roles back to the community as the more general purpose they are (and the wider the range of systems they will work on), the more useful they will be to people. Step through the following process to create your first role:\

1. Create the directory structure for the [installapache] role from within your chosen playbook directory---this is as simple as this:

```
$ mkdir -p roles/installapache/tasks
```

2. Now, let's create the mandatory [main.yml] inside the [tasks] directory we just created. This won't actually perform the Apache installation---rather, it will call one of two external tasks files, depending on the operating system detected on the target host during the fact-gathering stage. We can use this special variable, [ansible_distribution], in a [when] condition to determine which of the tasks files to import:

```
---
- name: import a tasks based on OS platform
  import_tasks: centos.yml
  when: ansible_distribution == 'CentOS'
- import_tasks: ubuntu.yml
  when: ansible_distribution == 'Ubuntu'
```

2. Create [centos.yml] in [roles/installapache/tasks] to install the latest version of the Apache web server via the [apt] package manager. This should contain the following content:

```
---
- name: Install Apache using apt
  apt:
    name: "apache2"
    state: latest
- name: Start the Apache server
  service:
    name: apache2
    state: started
```

3. Create a file called [ubuntu.yml] in [roles/installapache/tasks] to install the latest version of the Apache web server via the [apt] package manager on Ubuntu. Notice how the content differs between CentOS and Ubuntu hosts:

```
---
- name: Install Apache using apt
  apt:
    name: "apache2"
    state: latest
- name: Start the Apache server
  service:
    name: apache2
    state: started
```

Now, roles don't run by themselves---we have to create a playbook to call them, so let's write a simple playbook to call our newly created role. This has a play definition just like we saw before, but then rather than having a [tasks:] section within the play, we have a [roles:] section where the roles are declared instead. Convention dictates that this file be called [site.yml], but you are free to call it whatever you like:

```
---
- name: Install Apache using a role
  hosts: frontends
  become: true
```

```
roles:
  - installapache
```

For clarity, your final directory structure should look like this:

```
.
├── roles
│   ├── installapache
│   ├── tasks
│   ├── centos.yml
│   ├── main.yml
│   └── ubuntu.yml
└── site.yml
```

With this completed, you can now run your [site.yml] playbook using [ansible-playbook] in the normal way---you should see output similar to this:

```
$ cd ~/Desktop/gitlab-ci-ansible-course/Lab_16/role-example1
$ ansible-playbook -i hosts site.yml

PLAY [Install Apache using a role] *****

TASK [Gathering Facts] *****
ok: [frt01.example.com]
ok: [frt02.example.com]

TASK [installapache : Install Apache using apt] *****
changed: [frt02.example.com]
changed: [frt01.example.com]

TASK [installapache : Start the Apache server] *****
changed: [frt01.example.com]
changed: [frt02.example.com]

TASK [installapache : Install Apache using apt] *****
skipping: [frt01.example.com]
skipping: [frt02.example.com]

TASK [installapache : Start the Apache server] *****
skipping: [frt01.example.com]
skipping: [frt02.example.com]

PLAY RECAP *****
frt01.example.com : ok=3 changed=2 unreachable=0 failed=0 skipped=2 rescued=0
ignored=0
frt02.example.com : ok=3 changed=2 unreachable=0 failed=0 skipped=2 rescued=0
ignored=0
```

Since we have added hosts in `/etc/hosts` with `127.0.0.1` address so they are pointing to same machine, we might apt get locking error which can be ignored. This error won't occur using multiple machines

Before we look at some of the other aspects relating to roles, let's take a look at some other ways to call your role. Ansible allows you to statically import or dynamically include roles when you write a playbook. The syntax between

these importing or including a role is subtly different, and notably, both go in the tasks section of your playbook rather than in the roles section. The following is a hypothetical example that shows both options in a really simple playbook. The roles directory structure including both the [common] and [aprole] roles would have been created in a similar manner as in the preceding example:

```
---
- name: Play to import and include a role
  hosts: frontends

  tasks:
    - import_role:
        name: common
    - include_role:
        name: aprole
```

Setting up role-based variables and dependencies

Roles based variables can go in one of two locations:

- [defaults/main.yml]
- [vars/main.yml]

Note: Complete example is available at:

```
cd ~/Desktop/gitlab-ci-ansible-course/Lab_16/role-example2
```

We shall see examples of both of these as we build our example:

1. At the top level of our practical example, we will have a copy of the same inventory we have been using throughout this lab. We will also create a simple playbook called [site.yml], which contains the following code:

```
---
- name: Role variables and meta playbook
  hosts: frt01.example.com

  roles:
    - platform
```

Notice that we are simply calling one role called [platform] from this playbook---nothing else is called from the playbook itself.

2. Let's go ahead and create the [platform] role---unlike our previous role, this will not contain any tasks or even any variable data; instead, it will just contain a [meta] directory:

```
$ mkdir -p roles/platform/meta
```

Inside this directory, create a file called [main.yml] with the following contents:

```
---
dependencies:
- role: linuxtype
vars:
  type: centos
```

```
- role: linuxtype
  vars:
    type: ubuntu
```

This code will tell Ansible that the platform role is dependent on the [linuxtype] role. Notice that we are specifying the dependency twice, but each time we specify it, we are passing it a variable called [type] with a different value---in this way, the Ansible parser allows us to call the role twice because a different variable value has been passed to it each time it is referred to as a dependency.

3. Let's now go ahead and create the [linuxtype] role---again, this will contain no tasks, but more dependency declarations:

```
$ mkdir -p roles/linuxtype/meta/
```

Again, create a [main.yml] file in the [meta] directory, but this time containing the following:

```
---
dependencies:
- role: version
- role: network
```

Once again, we are creating more dependencies---this time, when the [linuxtype] role is called, it, in turn, is declaring dependencies on roles called [version] and [network].

4. Let's create the [version] role first---this will have both [meta] and [tasks] directories in it:

```
$ mkdir -p roles/version/meta
$ mkdir -p roles/version/tasks
```

In the [meta] directory, we'll create a [main.yml] file with the following contents:

```
---
allow_duplicates: true
```

This declaration is important in this example---as discussed earlier, normally Ansible will only allow a role to be executed once, even if it is called multiple times. Setting [allow_duplicates] to [true] tells Ansible to allow the execution of the role more than once. This is required because, in the [platform] role, we call (via a dependency) the [linuxtype] role twice, which means, in turn, we will call the [version] role twice.

We'll also create a simple [main.yml] file in the tasks directory, which prints the value of the [type] variable that gets passed to the role:

```
---
- name: Print type variable
  debug:
    var: type
```

5. We will now repeat the process with the [network] role---to keep our example code simple, we'll define it with the same contents as the [version] role:

```
$ mkdir -p roles/network/meta
$ mkdir -p roles/network/tasks
```

In the [meta] directory, we'll again create a [main.yml] file with the following contents:

```
---
allow_duplicates: true
```

Again, we'll create a simple [main.yml] file in the [tasks] directory, which prints the value of the [type] variable that gets passed to the role:

```
---
- name: Print type variable
  debug:
    var: type
```

At the end of this process, your directory structure should look like this:

```
.
├── hosts
├── roles
│   ├── linuxtype
│   │   ├── meta
│   │   └── main.yml
│   ├── network
│   │   ├── meta
│   │   │   └── main.yml
│   │   └── tasks
│   │       └── main.yml
│   ├── platform
│   │   ├── meta
│   │   └── main.yml
│   └── version
│       ├── meta
│       │   └── main.yml
│       └── tasks
│           └── main.yml
└── site.yml

11 directories, 8 files
```

Let's see what happens when we run this playbook. Now, you might think that the playbook is going to run like this: with the dependency structure we created in the preceding code, our initial playbook statically imports the [platform] role. The [platform] role then states that it depends upon the [linuxtype] role, and the dependency is declared twice with a different value in a variable called [type] each time. The [linuxtype] role then states that it depends upon both the [network] and [version] roles, which are allowed to run more than once and print the value of [type]. Hence, you could be forgiven for thinking that we'll see the [network] and [version] roles called twice, printing [centos] once and [ubuntu] the second time (as this is how we originally specified the dependencies in the [platform] role). However, when we run it, we actually see this:

```
$ ansible-playbook -i hosts site.yml

PLAY [Role variables and meta playbook] *****

TASK [Gathering Facts] *****
ok: [fvt01.example.com]

TASK [version : Print type variable] *****
```

```

ok: [frt01.example.com] => {
  "type": "ubuntu"
}

TASK [network : Print type variable] *****
ok: [frt01.example.com] => {
  "type": "ubuntu"
}

TASK [version : Print type variable] *****
ok: [frt01.example.com] => {
  "type": "ubuntu"
}

TASK [network : Print type variable] *****
ok: [frt01.example.com] => {
  "type": "ubuntu"
}

PLAY RECAP *****
frt01.example.com : ok=5 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0

```

If we are to continue using statically imported roles, then we should not use role variables when we declare the dependencies. Instead, we should pass over `[type]` as a role parameter. This is a small but crucial difference---role parameters remain scoped at the role level even when the Ansible parser is run, hence we can declare our dependency twice without the variable getting overwritten. To do this, change the contents of the `[roles/platform/meta/main.yml]` file to the following:

```

---
dependencies:
- role: linuxtype
  type: centos
- role: linuxtype
  type: ubuntu

```

Do you notice the subtle change? The `[vars:]` keyword has gone, and the declaration of `[type]` is now at a lower indentation level, meaning it is a role parameter. Now, when we run the playbook, we get the results that we had hoped for:

```

$ ansible-playbook -i hosts site.yml

PLAY [Role variables and meta playbook] *****

TASK [Gathering Facts] *****
ok: [frt01.example.com]

TASK [version : Print type variable] *****
ok: [frt01.example.com] => {
  "type": "centos"
}

TASK [network : Print type variable] *****

```

```

ok: [frr01.example.com] => {
    "type": "centos"
}

TASK [version : Print type variable] *****
ok: [frr01.example.com] => {
    "type": "ubuntu"
}

TASK [network : Print type variable] *****
ok: [frr01.example.com] => {
    "type": "ubuntu"
}

PLAY RECAP *****
frr01.example.com : ok=5 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0

```

This is quite an advanced example of Ansible role dependencies but it has been provided to you to demonstrate the importance of knowing a little about variable precedence (that is, where the variable is scoped) and how the parser works. If you write simple, sequentially parsed tasks, then you may never need to know this, but I recommend that you make extensive use of the debug statement and test your playbook design to make sure that you don't fall foul of this during your playbook development.

Having look in great detail at a number of aspects of roles, let's take a look in the following section at a centralized store for publicly available Ansible roles --- Ansible Galaxy.

Ansible Galaxy

No section on Ansible roles would be complete without a mention of Ansible Galaxy. Ansible Galaxy is a community-driven collection of Ansible roles, hosted by Ansible at <https://galaxy.ansible.com/>. It contains a great many community-contributed Ansible roles, and if you can conceive of an automation task, there is a good chance someone has already written a role to do exactly what you want it to do. It is well worth exploring and can get your automation project off the ground quickly as you can start work with a set of ready-made roles.

In addition to the web site, the [ansible-galaxy] client is included in Ansible, and this provides a quick and convenient way for you to download and deploy roles into your playbook structure. Let's say that you want to update the **message of the day (MOTD)** on your target hosts---this is surely something that somebody has already figured out. A quick search on the Ansible Galaxy website returns (at the time of writing) 106 roles for setting the MOTD. If we want to use one of these, we could download it into our roles directory using the following command:

```
$ ansible-galaxy role install -p roles/ arillso.motd
```

That's all you need to do---once the download is complete, you can import or include the role in your playbook just as you would for the manually created roles we have discussed in this lab. Note that if you don't specify [-p roles/], [ansible-galaxy] installs the roles into [~/./ansible/roles], the central roles directory for your user account. This might be what you want, of course, but if you want the role downloaded directly into your playbook directory structure, you would add this parameter.

Another neat trick is to use [ansible-galaxy] to create an empty role directory structure for you to create your own roles in---this saves all of the manual directory and file creation we have been undertaking in this lab, as in this example:


```
$ ansible-galaxy role init --init-path roles/ testrole
```

```
- Role testrole was created successfully
```

```
$ tree roles/testrole/
```

```
roles/testrole/
├── defaults
│   └── main.yml
├── files
├── handlers
│   └── main.yml
├── meta
│   └── main.yml
├── README.md
├── tasks
│   └── main.yml
├── templates
├── tests
│   ├── inventory
│   └── test.yml
└── vars
    └── main.yml
```

Using conditions in your code

In most of our examples so far, we have created simple sets of tasks that always run. However, as you generate tasks (whether in roles or playbooks) that you want to apply to a wider array of hosts, sooner or later, you will want to perform some kind of conditional action. This might be to only perform a task in response to the results of a previous task. Or it might be to only perform a task in response to a specific fact gathered from an Ansible system. In this section, we will provide some practical examples of conditional logic to apply to your Ansible tasks to demonstrate how to use this feature.

As ever, we'll need an inventory to get started, and we'll reuse the inventory we have used throughout this lab:

```
[frontends]
frt01.example.com https_port=8443
frt02.example.com http_proxy=proxy.example.com
```

```
[frontends:vars]
ntp_server=ntp.frt.example.com
proxy=proxy.frt.example.com
```

```
[apps]
app01.example.com
app02.example.com
```

```
[webapp:children]
frontends
apps
```

```
[webapp:vars]
proxy_server=proxy.webapp.example.com
```

```
health_check_retry=3
health_check_interval=60
```

Let's define the task that will perform our update but add a [when] clause containing a Jinja 2 expressions to it in a simple example playbook:

```
---
- name: Play to patch only CentOS systems
  hosts: all
  become: true

  tasks:
  - name: Patch Ubuntu systems
    apt:
      name: apache2
      state: latest
      when: ansible_facts['distribution'] == "Ubuntu"
```

Now, when we run this task, if your test system(s) are Ubuntu-based (and mine are), you should see output similar to the following:

```
$ cd ~/Desktop/gitlab-ci-ansible-course/Lab_16
$ ansible-playbook -i hosts condition.yml

PLAY [Play to patch only CentOS systems] *****

TASK [Gathering Facts] *****
ok: [frt02.example.com]
ok: [app01.example.com]
ok: [frt01.example.com]
ok: [app02.example.com]

TASK [Patch Ubuntu systems] *****
ok: [app01.example.com]
changed: [frt01.example.com]
ok: [app02.example.com]
ok: [frt02.example.com]

PLAY RECAP *****
app01.example.com : ok=2 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
app02.example.com : ok=2 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt01.example.com : ok=2 changed=1 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
frt02.example.com : ok=2 changed=0 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

The preceding output shows that all of our systems were Ubuntu-based, but that only [frt01.example.com] needed the patch applying. Now we can make our logic more precise---perhaps it is only our legacy systems that need the patch applying. In this case, we can expand the logic in our playbook to check both the distribution and major version, as follows:

```

---
- name: Play to patch only Ubuntu systems
  hosts: all
  become: true

  tasks:
  - name: Patch Ubuntu systems
    apt:
      name: apache2
      state: latest
    when: (ansible_facts['distribution'] == "Ubuntu" and
ansible_facts['distribution_major_version'] == "20")

```

Now, if we run our modified playbook, depending on the systems you have in your inventory, you might see output similar to the following. In this case, my [app01.example.com] server was based on CentOS 6 so had the patch applied. All other systems were skipped because they did not match my logical expression:

```

$ ansible-playbook -i hosts condition2.yml

PLAY [Play to patch only Ubuntu systems] *****

TASK [Gathering Facts] *****
ok: [frt01.example.com]
ok: [app02.example.com]
ok: [app01.example.com]
ok: [frt02.example.com]

TASK [Patch Ubuntu systems] *****
changed: [app01.example.com]
skipping: [frt01.example.com]
skipping: [frt02.example.com]
skipping: [app02.example.com]

PLAY RECAP *****
app01.example.com  : ok=2    changed=0    unreachable=0    failed=0    skipped=0
rescued=0    ignored=0
app02.example.com  : ok=2    changed=0    unreachable=0    failed=0    skipped=0
rescued=0    ignored=0
frt01.example.com  : ok=2    changed=0    unreachable=0    failed=0    skipped=0
rescued=0    ignored=0
frt02.example.com  : ok=2    changed=0    unreachable=0    failed=0    skipped=0
rescued=0    ignored=0

```

Consider the following playbook code. It contains two tasks, the first of which is to obtain the listing of the current directory and capture the output of the [shell] module in a variable called [shellresult]. When then print a simple [debug] message, but only on the condition that the [hosts] string is in the output of the [shell] command:

```

---
- name: Play to patch only CentOS systems
  hosts: localhost
  become: true

```

```

tasks:
  - name: Gather directory listing from local system
    shell: "ls -l"
    register: shellresult

  - name: Alert if we find a hosts file
    debug:
      msg: "Found hosts file!"
    when: '"hosts" in shellresult.stdout'

```

Now, when we run this in the current directory, which if you are working from the GitHub repository that accompanies this course will contain a file named [hosts], then you should see output similar to the following:

```

$ ansible-playbook condition3.yml

[WARNING]: provided hosts list is empty, only localhost is available. Note that
the implicit localhost does not match 'all'

PLAY [Play to patch only CentOS systems] *****

TASK [Gathering Facts] *****
ok: [localhost]

TASK [Gather directory listing from local system] *****
changed: [localhost]

TASK [Alert if we find a hosts file] *****
ok: [localhost] => {
  "msg": "Found hosts file!"
}

PLAY RECAP *****
localhost : ok=3 changed=1 unreachable=0 failed=0 skipped=0 rescued=0 ignored=0

```

Yet, if the file doesn't exist, then you'll see that the [debug] message gets skipped:

```

$ ansible-playbook condition3.yml

[WARNING]: provided hosts list is empty, only localhost is available. Note that
the implicit localhost does not match 'all'

PLAY [Play to patch only CentOS systems] *****

TASK [Gathering Facts] *****
ok: [localhost]

TASK [Gather directory listing from local system] *****
changed: [localhost]

TASK [Alert if we find a hosts file] *****
skipping: [localhost]

```

```
PLAY RECAP *****
localhost : ok=2 changed=1 unreachable=0 failed=0 skipped=1 rescued=0 ignored=0
```

You can also create complex conditions for IT operational tasks in production; however, remember that, in Ansible, variables are not cast to any particular type by default, and hence even though the contents of a variable (or fact) might look like a number, Ansible will by default treat it as a string. If you need to perform an integer comparison instead, you must first cast the variable to an integer type. For example, here is a fragment of a playbook that will run a task only on Fedora 25 and newer:

```
tasks:
  - name: Only perform this task on Fedora 25 and later
    shell: echo "only on Fedora 25 and later"
    when: ansible_facts['distribution'] == "Fedora" and
          ansible_facts['distribution_major_version']|int >= 25
```

There are many different types of conditionals you can apply to your Ansible tasks, and this section is just scratching the surface; however, it should give you a sound basis on which to expand your knowledge of applying conditions to your tasks in Ansible. Not only can you apply conditional logic to Ansible tasks, but you can also run them in loops over a set of data, and we shall explore this in the next section.

Repeating tasks with loops

Oftentimes, we will want to perform a single task, but use that single task to iterate over a set of data. For example, you might not want to create one user account but 10. Or you might want to install 15 packages to a system. The possibilities are endless, but the point remains the same---you would not want to write 10 individual Ansible tasks to create 10 user accounts. Fortunately, Ansible supports looping over datasets to ensure that you can perform large scale operations using tightly defined code. In this section, we will explore how to make practical use of loops in your Ansible playbooks.

As ever, we must start with an inventory to work against, and we will use our by-now familiar inventory, which we have consistently used throughout this lab:

```
[frontends]
frt01.example.com https_port=8443
frt02.example.com http_proxy=proxy.example.com

[frontends:vars]
ntp_server=ntp.frt.example.com
proxy=proxy.frt.example.com

[apps]
app01.example.com
app02.example.com

[webapp:children]
frontends
apps

[webapp:vars]
proxy_server=proxy.webapp.example.com
health_check_retry=3
health_check_interval=60
```

Let's start with a really simple playbook to show you how to loop over a set of data in a single task. Although this is quite a contrived example, it is intended to be simple to show you the fundamentals of how loops work in Ansible. We will define a single task that runs the [command] module on a single host from the inventory and uses the [command] module to [echo] the numbers 1 through 6 in turn on the remote system (with some imagination, this could easily be extended to adding user accounts or creating a sequence of files).

Consider the following code:

```
---
- name: Simple loop demo play
  hosts: frt01.example.com

  tasks:
    - name: Echo a value from the loop
      command: echo "{{ item }}"
      loop:
        - 1
        - 2
        - 3
        - 4
        - 5
        - 6
```

The [loop:] statement defines the start of the loop, and the items in the loop are defined as a YAML list. Also, note the higher indentation level, which tells the parser they are part of the loop. When working with the loop data, we use a special variable called [item], which contains the current value from the loop iteration to be echoed. Hence, if we run this playbook, we should see output similar to the following:

```
$ ansible-playbook -i hosts loop1.yml

PLAY [Simple loop demo play] *****

TASK [Gathering Facts] *****
ok: [frt01.example.com]

TASK [Echo a value from the loop] *****
changed: [frt01.example.com] => (item=1)
changed: [frt01.example.com] => (item=2)
changed: [frt01.example.com] => (item=3)
changed: [frt01.example.com] => (item=4)
changed: [frt01.example.com] => (item=5)
changed: [frt01.example.com] => (item=6)

PLAY RECAP *****
frt01.example.com : ok=2 changed=1 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0
```

You can combine the conditional logic we discussed in the preceding section with loops, to make the loop operate on just a subset of its data. For example, consider the following iteration of the playbook:

```
---
- name: Simple loop demo play
```

```

hosts: frt01.example.com

tasks:
  - name: Echo a value from the loop
    command: echo "{{ item }}"
    loop:
      - 1
      - 2
      - 3
      - 4
      - 5
      - 6
    when: item|int > 3

```

Now, when we run this, we see that the task is skipped until we reach the integer value of 4 and higher in the loop contents:

```

$ ansible-playbook -i hosts loop2.yml

PLAY [Simple loop demo play] *****

TASK [Gathering Facts] *****
ok: [frt01.example.com]

TASK [Echo a value from the loop] *****
skipping: [frt01.example.com] => (item=1)
skipping: [frt01.example.com] => (item=2)
skipping: [frt01.example.com] => (item=3)
changed: [frt01.example.com] => (item=4)
changed: [frt01.example.com] => (item=5)
changed: [frt01.example.com] => (item=6)

PLAY RECAP *****
frt01.example.com : ok=2 changed=1 unreachable=0 failed=0 skipped=0 rescued=0
ignored=0

```

You can, of course, combine this with the conditional logic based on Ansible facts and other variables in the manner we discussed previously. Just as we captured the results of a module's execution using the `[register]` keyword before, we can do so with loops. The only difference is that the results will now be stored in a dictionary, with one dictionary entry for each iteration of the loop rather than just one set of results.

Hence, let's see what happens if we further enhance the playbook, as follows:

```

---
- name: Simple loop demo play
  hosts: frt01.example.com

  tasks:
    - name: Echo a value from the loop
      command: echo "{{ item }}"
      loop:
        - 1
        - 2

```

```

- 3
- 4
- 5
- 6
when: item|int > 3
register: loopresult

- name: Print the results from the loop
  debug:
    var: loopresult

```

Now, when we run the playbook, you will see pages out output containing the dictionary with the contents of [loopresult]. The following output is truncated in the interests of space but demonstrates the kind of results you should expect from running this playbook:

```

$ ansible-playbook -i hosts loop3.yml

PLAY [Simple loop demo play] *****

TASK [Gathering Facts] *****
ok: [frr01.example.com]

TASK [Echo a value from the loop] *****
skipping: [frr01.example.com] => (item=1)
skipping: [frr01.example.com] => (item=2)
skipping: [frr01.example.com] => (item=3)
changed: [frr01.example.com] => (item=4)
changed: [frr01.example.com] => (item=5)
changed: [frr01.example.com] => (item=6)

TASK [Print the results from the loop] *****
ok: [frr01.example.com] => {
  "loopresult": {
    "changed": true,
    "msg": "All items completed",
    "results": [
      {
        "ansible_loop_var": "item",
        "changed": false,
        "item": 1,
        "skip_reason": "Conditional result was False",
        "skipped": true
      },
      {
        "ansible_loop_var": "item",
        "changed": false,
        "item": 2,
        "skip_reason": "Conditional result was False",
        "skipped": true
      },

```

As you can see, the results section of the output is a dictionary, and we can clearly see that the first two items in the list were [skipped] because the result of our [when] clause ([Conditional]) was [false].

First of all, we'll create a playbook in the usual manner, with a single task to run in a loop. To generate our nested loop, we'll use the `[include_tasks]` directory to dynamically include a single task from another YAML file that will also contain a loop. As we're intending to use this playbook in a nested loop, we'll use the `[loop_var]` directive to change the name of the special loop contents variable from `[item]` to `[second_item]`:

```
---
- name: Play to demonstrate nested loops
  hosts: localhost

  tasks:
    - name: Outer loop
      include_tasks: loopsubtask.yml
      loop:
        - a
        - b
        - c
      loop_control:
        loop_var: second_item
```

Then, we'll create a second file called `[loopsubtask.yml]`, which contains the inner loop and is included in the preceding playbook. As we're already changed the loop item variable name in the outer loop, we don't need to change it again here. Note that the structure of this file is very much like a tasks file in a role---it is not a complete playbook, but rather simply a list of tasks:

```
---
- name: Inner loop
  debug:
    msg: "second item={{ second_item }} first item={{ item }}"
  loop:
    - 100
    - 200
    - 300
```

Now you should be able to run the playbook, and you will see Ansible iterate over the outer loop first and then process the inner loop over the data defined by the outer loop. As the loop variable names do not clash, all works exactly as we would expect:

```
$ ansible-playbook loopmain.yml

PLAY [Play to demonstrate nested loops] *****

TASK [Gathering Facts] *****
ok: [localhost]

TASK [Outer loop] *****
included: /root/Desktop/gitlab-ci-ansible-course/Lab_16/loopsubtask.yml for localhost
included: /root/Desktop/gitlab-ci-ansible-course/Lab_16/loopsubtask.yml for localhost
included: /root/Desktop/gitlab-ci-ansible-course/Lab_16/loopsubtask.yml for localhost

TASK [Inner loop] *****
ok: [localhost] => (item=100) => {
```

```

    "msg": "second item=a first item=100"
  }
ok: [localhost] => (item=200) => {
    "msg": "second item=a first item=200"
  }
ok: [localhost] => (item=300) => {
    "msg": "second item=a first item=300"
  }

TASK [Inner loop] *****
ok: [localhost] => (item=100) => {
    "msg": "second item=b first item=100"
  }
ok: [localhost] => (item=200) => {
    "msg": "second item=b first item=200"
  }
ok: [localhost] => (item=300) => {
    "msg": "second item=b first item=300"
  }

TASK [Inner loop] *****
ok: [localhost] => (item=100) => {
    "msg": "second item=c first item=100"
  }
ok: [localhost] => (item=200) => {
    "msg": "second item=c first item=200"
  }
ok: [localhost] => (item=300) => {
    "msg": "second item=c first item=300"
  }

PLAY RECAP *****
localhost : ok=7 changed=0 unreachable=0 failed=0 skipped=0 rescued=0 ignored=0

```

Loops are simple to work with, and yet very powerful as they allow you to easily use one task to iterate over a large set of data. In the next section, we'll look at another construct of the Ansible language for controlling playbook flow--blocks.

Grouping tasks using blocks

Blocks in Ansible allow you to logically group a set of tasks together, primarily for one of two purposes. One might be to apply conditional logic to an entire set of tasks; in this example, you could apply an identical `when` clause to each of the tasks, but this is cumbersome and inefficient---far better to place all of the tasks in a block and apply the conditional logic to the block itself. In this way, the logic only needs to be declared once. Blocks are also valuable when it comes to error handling and especially when it comes to recovering from an error condition. We shall explore both of these through simple practical examples in this lab to get you up to speed with blocks in Ansible.

As ever, let's ensure we have an inventory to work from:

```

[frontends]
frt01.example.com https_port=8443
frt02.example.com http_proxy=proxy.example.com

```

```

[frontends:vars]
ntp_server=ntp.frt.example.com
proxy=proxy.frt.example.com

[apps]
app01.example.com
app02.example.com

[webapp:children]
frontends
apps

[webapp:vars]
proxy_server=proxy.webapp.example.com
health_check_retry=3
health_check_interval=60

```

Now, let's dive straight in and look at an example of how you would use blocks to apply conditional logic to a set of tasks. At a high level, suppose we want to perform the following actions on all of our Fedora Linux hosts:

- Install the package for the Apache web server.
- Install a templated configuration.
- Start the appropriate service.

We could achieve this with three individual tasks, all with a [when] clause associated with them, but blocks provide us with a better way. The following example playbook shows the three tasks discussed contained in a block (notice the additional level of indentation required to denote their presence in the block):

```

---
- name: Conditional block play
  hosts: all
  become: true

  tasks:
    - name: Install and configure Apache
      block:
        - name: Install the Apache package
          dnf:
            name: apache2
            state: installed
        - name: Install the templated configuration to a dummy location
          template:
            src: templates/src.j2
            dest: /tmp/my.conf
        - name: Start the apache2 service
          service:
            name: apache2
            state: started
            enabled: True
      when: ansible_facts['distribution'] == 'Fedora'

```

When you run this playbook, you should find that the Apache-related tasks are only run on any Fedora hosts you might have in your inventory; you should see that either all three tasks are run or are skipped---depending on the makeup and contents of your inventory, it might look something like this:

```
$ ansible-playbook -i hosts blocks.yml

PLAY [Conditional block play] *****

TASK [Gathering Facts] *****
ok: [app02.example.com]
ok: [frt01.example.com]
ok: [app01.example.com]
ok: [frt02.example.com]

TASK [Install the Apache package] *****
changed: [frt01.example.com]
changed: [frt02.example.com]
skipping: [app01.example.com]
skipping: [app02.example.com]

TASK [Install the templated configuration to a dummy location] *****
changed: [frt01.example.com]
changed: [frt02.example.com]
skipping: [app01.example.com]
skipping: [app02.example.com]

TASK [Start the apache2 service] *****
changed: [frt01.example.com]
changed: [frt02.example.com]
skipping: [app01.example.com]
skipping: [app02.example.com]

PLAY RECAP *****
app01.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=3 rescued=0
ignored=0
app02.example.com : ok=1 changed=0 unreachable=0 failed=0 skipped=3 rescued=0
ignored=0
frt01.example.com : ok=4 changed=3 unreachable=0 failed=0 skipped=3 rescued=0
ignored=0
frt02.example.com : ok=4 changed=3 unreachable=0 failed=0 skipped=3 rescued=0
ignored=0
```

This is very simple to construct, but very powerful in terms of the effect it has on your ability to control the flow over large sets of tasks.

This time, let's build a different example to demonstrate how blocks can be utilized to help Ansible to handle error conditions gracefully. So far, you should have experienced that if your playbooks encounter any errors, they are likely to stop executing at the point of failure. This is in some situations far from ideal, and you might want to perform some kind of recovery actions in this event rather than simply halting the playbook.

Let's create a new playbook, this time with the following contents:

```

---
- name: Play to demonstrate block error handling
  hosts: frontends

  tasks:
    - name: block to handle errors
      block:
        - name: Perform a successful task
          debug:
            msg: 'Normally executing....'
        - name: Deliberately create an error
          command: /bin/whatever
        - name: This task should not run if the previous one results in an error
          debug:
            msg: 'Never print this message if the above command fails!!!!'
      rescue:
        - name: Catch the error (and perform recovery actions)
          debug:
            msg: 'Caught the error'
        - name: Deliberately create another error
          command: /bin/whatever
        - name: This task should not run if the previous one results in an error
          debug:
            msg: 'Do not print this message if the above command fails!!!!'
      always:
        - name: This task always runs!
          debug:
            msg: "Tasks in this part of the play will be ALWAYS executed!!!"

```

Notice that in the preceding play, we now have additional sections to [block] --- as well as the tasks in [block] itself, we have two new parts labeled [rescue] and [always]. The flow of execution is as follows:

With this flow of execution in mind, you should see output similar to the following when you execute this playbook, noting that we have deliberately created two error conditions to demonstrate the flow:

```

$ ansible-playbook -i hosts blocks-error.yml

PLAY [Play to demonstrate block error handling] *****

TASK [Gathering Facts] *****
ok: [frt02.example.com]
ok: [frt01.example.com]

TASK [Perform a successful task] *****
ok: [frt01.example.com] => {
  "msg": "Normally executing...."
}
ok: [frt02.example.com] => {
  "msg": "Normally executing...."
}

TASK [Deliberately create an error] *****

```

```

fatal: [frt01.example.com]: FAILED! => {"changed": false, "cmd": "/bin/whatever",
"msg": "[Errno 2] No such file or directory", "rc": 2}
fatal: [frt02.example.com]: FAILED! => {"changed": false, "cmd": "/bin/whatever",
"msg": "[Errno 2] No such file or directory", "rc": 2}

TASK [Catch the error (and perform recovery actions)] *****
ok: [frt01.example.com] => {
    "msg": "Caught the error"
}
ok: [frt02.example.com] => {
    "msg": "Caught the error"
}

TASK [Deliberately create another error] *****
fatal: [frt01.example.com]: FAILED! => {"changed": false, "cmd": "/bin/whatever",
"msg": "[Errno 2] No such file or directory", "rc": 2}
fatal: [frt02.example.com]: FAILED! => {"changed": false, "cmd": "/bin/whatever",
"msg": "[Errno 2] No such file or directory", "rc": 2}

TASK [This task always runs!] *****
ok: [frt01.example.com] => {
    "msg": "Tasks in this part of the play will be ALWAYS executed!!!"
}
ok: [frt02.example.com] => {
    "msg": "Tasks in this part of the play will be ALWAYS executed!!!"
}

PLAY RECAP *****
frt01.example.com : ok=4 changed=0 unreachable=0 failed=1 skipped=0 rescued=1
ignored=0
frt02.example.com : ok=4 changed=0 unreachable=0 failed=1 skipped=0 rescued=1
ignored=0

```

Configuring play execution via strategies

The most useful strategy when you are starting work with Ansible, however, is going to be the `[debug]` strategy, and this enables Ansible to drop you straight into its integrated debug environment if an error should occur in the playbook. Let's demonstrate this by creating a playbook that has a deliberate error in it. Note the `[strategy: debug]` and `[debugger: on_failed]` statements in the play definition:

```

---
- name: Play to demonstrate the debug strategy
  hosts: frt01.example.com
  strategy: debug
  debugger: on_failed
  gather_facts: no
  vars:
    username: daniel

  tasks:

```

```
- name: Generate an error by referencing an undefined variable
ping: data={{ mobile }}
```

Now if you execute this playbook, you should see that it starts to run, but then drops you into the integrated debugger when it encounters the deliberate error it contains. The start of the output should be similar to the following:

```
$ ansible-playbook -i hosts debug.yml

PLAY [Play to demonstrate the debug strategy] *****

TASK [Generate an error by referencing an undefined variable] *****
fatal: [ftr01.example.com]: FAILED! => {"msg": "The task includes an option with an
undefined variable. The error was: 'mobile' is undefined\n\nThe error appears to be in
'/root/Desktop/gitlab-ci-ansible-course/Lab_16/debug.yml': line 11, column 7, but
may\nbe elsewhere in the file depending on the exact syntax problem.\n\nThe offending
line appears to be:\n\n tasks:\n - name: Generate an error by referencing an undefined
variable\n ^ here\n"}
[ftr01.example.com] TASK: Generate an error by referencing an undefined variable
(debug)>

[ftr02.prod.com] TASK: make an error with refering incorrect variable (debug)> p
task_vars
{'ansible_check_mode': False,
 'ansible_current_hosts': [u'ftr02.prod.com'],
 'ansible_diff_mode': False,
 'ansible_facts': {},
 'ansible_failed_hosts': [],
 'ansible_forks': 5,
 ...
[ftr02.prod.com] TASK: make an error with refering incorrect variable (debug)> quit
User interrupted execution
$
```

To take you through a very simple, practical debugging example, however, enter the `[p task]` command at the prompt---this will cause the Ansible debugger to print the name of the failing task; this is very useful if you are in the midst of a large playbook:

```
[ftr01.example.com] TASK: Generate an error by referencing an undefined variable
(debug)> p task
TASK: Generate an error by referencing an undefined variable
```

Now we know where the play failed, so let's dig a little deeper by issuing the `[p task.args]` command, which will show us the arguments that were passed to the module in the task:

```
[ftr01.example.com] TASK: Generate an error by referencing an undefined variable
(debug)> p task.args
{'data': u'{{ mobile }}'}
```

So, we can see that our module was passed the argument called `[data]`, with the argument value being a variable (denoted by the pairs of curly braces) called `[mobile]`. Hence, it might be logical to have a look at the variables

available to the task, to see whether this variable exists, and if so whether the value is sensible (use the [p task_vars] command to do this):

```
[frt01.example.com] TASK: Generate an error by referencing an undefined variable
(debug)> p task_vars
{'ansible_check_mode': False,
 'ansible_current_hosts': [u'frt01.example.com'],
 'ansible_dependent_role_names': [],
 'ansible_diff_mode': False,
 'ansible_facts': {},
 'ansible_failed_hosts': [],
 'ansible_forks': 5,
```

The preceding output is truncated, and you will find a great many variables associated with the task---this is because any gathered facts, and internal Ansible variables, are all available to the task. However, if you scroll through the list, you will be able to confirm that there is no variable called [mobile].

Hence, this should be enough information to fix your playbook. Enter [q] to quit the debugger:

```
[frt01.example.com] TASK: Generate an error by referencing an undefined variable
(debug)> q
User interrupted execution
$
```

The Ansible debugger is an incredibly powerful tool and you should learn to make effective use of it, especially as your playbook complexity grows. This concludes our practical look at the various aspects of playbook design---in the next section, we'll take a look at the ways in which you can integrate Git source code management into your playbooks.

Using ansible-pull

Let's use a simple playbook from GitHub that sets the message of the day based on variable content. To do this, we will run the following command (which we'll break down in a minute):

```
$ ansible-pull -d /var/ansible-set-motd -i ${HOSTNAME}, -U
https://github.com/fenago/ansible-set-motd.git site.yml -e "ag_motd_content='MOTD
generated by ansible-pull'" >> /tmp/ansible-pull.log 2>&1
```

When you run this command, you should see some output similar to the following (note that log redirection has been removed to make it easier to see the output):

```
$ ansible-pull -d /var/ansible-set-motd -i ${HOSTNAME}, -U
https://github.com/fenago/ansible-set-motd.git site.yml -e "ag_motd_content='MOTD
generated by ansible-pull'"

Starting Ansible Pull at 2020-04-14 17:26:21
/usr/bin/ansible-pull -d /var/ansible-set-motd -i cookbook, -U
https://github.com/fenago/ansible-set-motd.git site.yml -e ag_motd_content='MOTD
generated by ansible-pull'
```



```

cookbook |[WARNING]: SUCCESS = Your git > {
  "aversion isfter": "7d too old t3a191ecb2do fully suebe7f84f4fpport the
a5817b0f1bdepth argu49c4cd54",ment.
Fall
  "ansing back tible_factso full che": {
    ckouts.
  "discovered_interpreter_python": "/usr/bin/python"
  },
  "before": "7d3a191ecb2debe7f84f4fa5817b0f1b49c4cd54",
  "changed": false,
  "remote_url_changed": false
}

PLAY [Update the MOTD on hosts] *****

TASK [Gathering Facts] *****
ok: [cookbook]

TASK [ansible.motd : Add 99-footer file] *****
skipping: [cookbook]

TASK [ansible.motd : Delete 99-footer file] *****
ok: [cookbook]

TASK [ansible.motd : Delete /etc/motd file] *****
skipping: [cookbook]

TASK [ansible.motd : Check motd tail supported] *****
fatal: [cookbook]: FAILED! => {"changed": true, "cmd": "test -f /etc/update-motd.d/99-
footer", "delta": "0:00:00.004444", "end": "2020-04-14 17:26:25.489793", "msg": "non-
zero return code", "rc": 1, "start": "2020-04-14 17:26:25.485349", "stderr": "",
"stderr_lines": [], "stdout": "", "stdout_lines": []}
...ignoring

TASK [ansible.motd : Add motd tail] *****
skipping: [cookbook]

TASK [ansible.motd : Add motd] *****
changed: [cookbook]

PLAY RECAP *****
cookbook : ok=4 changed=2 unreachable=0 failed=0 skipped=3 rescued=0 ignored=1

```

This command can be a very powerful part of your overall Ansible solution, especially as it means you don't have to worry too greatly about running all of your playbooks centrally, or ensuring that they are all up to date every time you run them. The ability to schedule this in [cron] is especially powerful in a large infrastructure where, ideally, automation means things should take care of themselves.

This concludes our practical look at playbooks and how to author your own code---with a little research into Ansible modules, you should now have enough to write your own robust playbooks with ease.

Summary

In this lab, you learned about the playbook framework and how to start building your own playbooks. You then learned how to organize your code into roles and design your code to effectively and efficiently support reuse. We then explored some of the more advanced playbook writing topics such as working with conditional logic, blocks, and loops. Finally, we looked at playbook execution strategies, especially with a view to being able to debug your playbooks effectively, and we wrapped up with a look at how you can run Ansible playbooks on a local machine directly from GitHub.

In the next lab, we will learn how to consume and create our very own modules, providing you with the skills you need to expand the capabilities of Ansible to suit your own bespoke environments, and to contribute back to the community.

Questions

1. How do you restart the Apache web server in the [frontends] host group via an ad hoc command?

- A) `[ansible frontends -i hosts -a "name=apache2 state=restarted"]`
- B) `[ansible frontends -i hosts -b service -a "name=apache2 state=restarted"]`
- C) `[ansible frontends -i hosts -b -m service -a "name=apache2 state=restarted"]`
- D) `[ansible frontends -i hosts -b -m server -a "name=apache2 state=restarted"]`
- E) `[ansible frontends -i hosts -m restart -a "name=apache2"]`

2. Do blocks allow you to logically make a group of tasks, or perform error handling?

- A) True
- B) False

3. Default strategies are intended via the relevant modules in the playbook.

- A) True
- B) False